



Matthias Köppe <matthiaskoepp@gmail.com>

Sage Code of Conduct reply to 'Examples of "acting offended"'

2 messages

Nils Bruin <bruin.nils@gmail.com>

Fri, May 3, 2024 at 9:42 AM

To: Matthias Koepp <mkoeppe@math.ucdavis.edu>, sage-conduct@googlegroups.com

Dear Matthias,

This is a reply from the CoCC to several messages you sent to us illustrating certain behaviours you attribute to Gonzalo Tornaria. You assess that he is “acting offended” as a performative manipulation. While we cannot read Gonzalo’s mind and hence cannot with certainty determine his motivations, we have reviewed the referenced interactions and we have come to the conclusion that it is very credible that Gonzalo is not acting, but is expressing his genuine reactions and emotions that are triggered by your responses.

The behavioural patterns we observe on your side are consistent with the patterns we mentioned in our message *Sage Code of Conduct Committee response and assessment regarding several reports* (sent April 15 2024) and we think they are likely to trigger offense, frustration, and irritation in most humans. The fact that you are surprised by the effect may indicate that you are not aware of how these patterns of communication come across in discussion forums. Electronic, written communication is notoriously prone to misunderstanding because it misses many of the nonverbal cues present in in-person communication that express tone and intent.

Below we will give our analysis of the discussion fragments you pointed to and explain how it reads to a bystander and therefore very likely also to your discussion partner. These case studies will hopefully help you to recognize these types of communication patterns in the future when they occur and perhaps avoid them.

Fragment 1 (https://github.com/sagemath/sage/pull/37785#discussion_r1564787139):

tornaria: "don't tell me what I should do with my time" (I sure didn't)

Analysis:

The context here is that on <https://github.com/sagemath/sage/pull/36561>, particularly in the last commit <https://github.com/sagemath/sage/pull/37391/commits/cb4e7184b3af09a1ed3c2353aba77d63f206276d>, a non-portable option “time -h” is introduced, instead of the more widely available option “time -p”. The PR here deals with a diagnosed extraneous error message if `/usr/bin/time`` is not available, by hiding the stderr. On linux, the line would fail even if `/usr/bin/time` is available, but now silently, leading to the variable “\$TIME” not being configured. That this change made it through testing and merging is a break-down in the review process of sage.

It is quite credible that Gonzalo is expressing genuine frustration here. It is quite believable that he felt the issue he was pointing out, that “time -h” isn’t portable, was not being acknowledged. He gives a whole bunch of very constructive suggestions. These are not acknowledged at the point of his comment at all. Instead there is just a “The whole thing is just a convenience.” and a neutral “If you want to prepare a follow-up PR, I can test it for its portability. I don't have time to work on it myself at the moment.”. It is quite believable that at this point Gonzalo read this as “if you know it so well, then do it yourself”. He actually

suggests an alternative course of action, just revert the use of `time -h` back to `time -p`, that does not require any development, punting to a time in the future to get a portable alternative for producing human-readable times, when someone does have the time.

An alternative way to formulate this would have been: “indeed, those are good suggestions for a portable alternative for “time -h”, but I myself don’t have time right now to pursue those and I think the problem on this ticket is urgent enough that it’s beneficial to put this temporary fix in here. Reverting to “time -p” does not work because

There now is a follow-up PR <https://github.com/sagemath/sage/pull/37785> where this work is happening, but the dialogue there reads like the distrust sowed by the disagreement here have made all parties much more skeptical about solutions offered by the others, slowing down discussions there as well. Investing some time in acknowledging contributions and comments from community members will probably pay off in future collaborations.

Fragment 2 (<https://github.com/sagemath/sage/pull/37287#issuecomment-2067778819>)

mkoepppe: [@tornaria](#) That's false. It's testable using just sagemath-repl, which provides the doctester.

tornaria: "This sounds a little bit aggressive. We are both making an effort to communicate in better terms, so let's try to keep it cool and learn to agree and disagree.

You know what I mean. I'm not currently packaging sagemath-repl since that way of packaging is not ready and not working and, AFAIK, sagemath-repl and sagemath-standard *do* overlap in the binary distributions. Am I mistaken?"

Analysis:

It’s a little unclear what Gonzalo means by “aggressive” here. However, he does follow that opening sentence with a further clarification:

“You know what I mean. I'm not currently packaging sagemath-repl since that way of packaging is not ready and not working and, AFAIK, sagemath-repl and sagemath-standard *do* overlap in the binary distributions. Am I mistaken?"

He is referring to your statement: “It's testable using just **sagemath-repl**, which provides the doctester.”

Gonzalo flags what he perceives as a miscommunication here: he thinks your answer doesn’t address the point he was trying to make and he points out why that is the case.

Note that statements like “That’s false” can come across as harsh. It looks like Gonzalo felt an unfriendly undertone in the message. He is referring to a history of less-than-optimal communication, so that could be a reason why he was more prone to reading negative connotations in a message. He is reaching out with “we are both making an effort”. It is probably more productive to concentrate on the conciliatory part and not dwell too much on a somewhat poorly formulated phrase.

Fragment 3 (<https://github.com/sagemath/sage/pull/35404> , Drop support for Python 3.8):

mkoepppe: “Tobias, over in [#35403 \(comment\)](#), you acknowledged having seen my comments in which I explain what factors into a decision to drop support for a Python version.”

tobiasdiez: “Yes, I've seen your comments concerning NEP29. They are not relevant here as this PR is independent of NEP29.

Above you say also: 'We will have a strong case for dropping support for Python 3.8 when we want to upgrade major packages that have done this.'

But the drop schedule of other packages is not relevant here. And anyway, it's not saying anything about why Python 3.8 support shouldn't be dropped now."

mkoeppe: "Tobias, you are too quick in dismissing my points as "not relevant". I think you should take the time to think about this more."

tobiasdiez: "Do you really think ad hominem attacks are helpful?"

From what I can see dependencies needing a higher python version can serve as an argument why sage should drop a certain version. But this is certainly not a necessary (nor sufficient) criteria. In particular, I fail to see how the fact that scipy dropped Python 3.8 in a PR this February implies that we cannot drop the same version in a PR in April."

Analysis: Agreed that the label "ad hominem attack" is not an accurate description of what is written. The message "I think you should take the time to think about this more" does come across as patronizing because it dismisses the argument made by the other party as poorly thought out rather than pointing out concretely what you don't agree with in the rebuttal: in this case, Tobias says that he has considered the comments on NEP29 and has come to the conclusion that those are not relevant. What you're saying can be read as "yes they are. I leave it as a homework exercise for you to figure out why". Tobias could very well have felt personally attacked there and reacted to that rather than what was literally written.

Quite frankly, it's not clear the discussion on this ticket could be rescued at this point anyway: two parties have clearly dug in on their positions and at this point the discussion has devolved into a debate where the two parties only have as a goal to convince the other side that they are right. Those goals are fundamentally in conflict with each other, so resolution by continuing the discussion is not going to happen anyway. That's a flaw in the process of getting tickets accepted into sage. That's what led to the "disputed ticket" procedure, which has its own flaws.

It is understandable that conversation partners in such debate situations reach to not-friendly language, because the conversation at that point does have strong adversarial and competitive elements. By engaging in it you will reduce trust and goodwill, which will make collaborations on later issues more difficult. This includes the trust and goodwill of bystanders. Avoid patronizing language and blanket dismissals if possible.

Fragment 4 (https://github.com/sagemath/sage/pull/36561#discussion_r1575449665):

tornaria: "Promote out of beta" doesn't mean it didn't work before 68.1.0, but as I said above, it's not that important.

ISTM that 61.0 might be good enough, but again, this is not so important.

Cf. <https://setuptools.pypa.io/en/latest/history.html#v61-0-0>

mkoeppe: Surely you are not suggesting that it would be an improvement to use beta-quality features for our production work.

tornaria: I don't know if you are trying to play some kind of game here. I rather play with my kids. Here I'm trying to help you merge this PR.

mkoeppe: I am not; and this is an unwelcome speculation.

tornaria: Do you realize I've felt unwelcomed way before your last comment?

Analysis: The PR is about reorganizing metadata and in the process ups the minimum version requirement for setuptools to a relatively new one (68.1.1). Tornaria points out various drawbacks of this, among which that (at the time) several popular linux distributions do not ship such a recent setuptools. The reason for the version lower bound given is that with 68.1.0, the required option was declared out of beta. However, as Tornaria comments, the required functionality was available long before. Indeed the changelogs of setuptools do not indicate any significant changes for a long time. From that perspective, dismissing a compromise on such formal grounds may well look like an overly bureaucratic ploy to avoid making any changes. It is quite probable that Tornaria felt he pointed out that the required features had been stable for a long while to you in his earlier comment and hence was irritated that something he considered was a valuable piece of information was dismissed on bureaucratic grounds. From this point it's clear that the discussion will not lead to consensus anyway and essentially Tornaria signs off in desperation with a snarky message, reiterating a reason why he thought the discussion would be attractive for you to engage in seriously. You follow up with a denial, which further widens the gap between the parties.

It does not look like Tornaria is "playing the victim" in expressing his feeling of being unwelcome. In fact, for a bystander reading through the exchange, Tornaria raises various credible concerns together with fairly simple solutions which are all met with outright dismissal, even though he sketches reasons and incentives for you to engage in reaching a compromise. It does read like his comments were unwelcome from the start.

Interestingly, the pull request opens with:

This is a trivial "chore" PR. It updates Python metadata to the latest format. No controversies about the current format are known about in the Python community. In a typical open source project, someone in a Maintainer role would open a PR and then immediately merge it

Whereas the comments make clear it turns out to be not so trivial after all: the version requirement bump is a non-trivial change and has real knock-on effects. Designating a ticket as "trivial" can undermine trust and put more friction on future interactions and collaborations.

The Code of Conduct Committee

Matthias Koeppe <mkoeppe@math.ucdavis.edu>
To: Nils Bruin <bruin.nils@gmail.com>
Cc: sage-conduct@googlegroups.com

Sat, May 4, 2024 at 9:56 PM

Dear Nils and committee,

On Fri, May 3, 2024 at 9:43AM Nils Bruin <bruin.nils@gmail.com> wrote:

> This is a reply from the CoCC to several messages you sent to us illustrating certain behaviours you attribute to Gonzalo Tornaria.

> You assess that he is "acting offended" as a performative manipulation. While we cannot read Gonzalo's mind and hence cannot with certainty determine his motivations [...]

That's right, you can't!

> Below we will give our analysis of the discussion fragments you pointed to and explain how it reads to a bystander

Thanks for carefully working on my prompts. That's a good contribution to our discussion. Indeed, how the discussion reads to a bystander is very relevant! Let's take it as a starting point, and I'll suggest the next steps in the analysis below.

> and therefore very likely also to your discussion partner.

Here you are already _assuming_ the good faith in that "discussion

partner" and ignoring/denying the very relevant possibility of the use of manipulative techniques.

I also note that you are saying "very likely" here.

I agree with you that our community has only a very small number of individuals who engage in manipulative and abusive conduct.

Hence I agree that if we were to look at a discussion sampled uniformly from the corpus, then it would be "very likely" that how a discussion partner reads a comment will align with how a 'bystander' reads it.

But the discussions that are brought to the attention of the committee are obviously not the result of uniform sampling; so just as a general point, I would suggest caution in using expressions of probability.

But more to the specific point, manipulation succeeds in part _because_ it uses techniques that work and in part _because_ it is rare.

For the committee to deny manipulation on the basis that (1) a bystander would not see through it and (2) that it's "unlikely" --- that's fundamentally flawed; and it's necessary to work on this.

> These case studies will hopefully help you to recognize these types of communication patterns in the future when they occur and perhaps avoid them.

Thank you, but I think you are being a bit quick here, and moreover I get a certain sense of receiving advice to wear a longer skirt.

But let's go to the examples (I've made a selection to keep the post shorter).

> Fragment 2 (<https://github.com/sagemath/sage/pull/37287#issuecomment-2067778819>)

> mkoeppe: @tornaria That's false. It's testable using just sagemath-repl, which provides the doctester.

> tornaria: "This sounds a little bit aggressive. We are both making an effort to communicate in better terms, so let's try to keep it cool and learn to agree and disagree."

> You know what I mean. I'm not currently packaging sagemath-repl since that way of packaging is not ready and not working and, AFAIK, sagemath-repl and sagemath-standard do overlap in the binary distributions. Am I mistaken?"

>

> Analysis:

>

> It's a little unclear what Gonzalo means by "aggressive" here.

It's not unclear at all. He is playing offended for the audience; it's a manipulation.

That you say it's "a little unclear" means that he has succeeded in creating plausible deniability -- a key component of a successful manipulation.

> [...] He is reaching out with "we are both making an effort".

No, he's not "reaching out". He's the abuser and with this "both of us" he is seeking to manipulate the audience into perceiving a symmetric situation. (It seems his manipulation may have succeeded with you.)

> Fragment 3 (<https://github.com/sagemath/sage/pull/35404> , Drop support for Python 3.8):

> mkoeppe: "Tobias, over in #35403 (comment), you acknowledged having seen my comments in which I explain what factors into a decision to drop support for a Python version."

> tobiasdiaz: "Yes, I've seen your comments concerning NEP29. They are not relevant here as this PR is independent of NEP29."

> Above you say also: 'We will have a strong case for dropping support for Python 3.8 when we want to upgrade major packages that have done this.'

> But the drop schedule of other packages is not relevant here. And anyway, it's not saying anything about why Python 3.8 support shouldn't be dropped now."

> mkoepe: "Tobias, you are too quick in dismissing my points as "not relevant". I think you should take the time to think about this more."
> tobiasdiez: "Do you really think ad hominem attacks are helpful? [...]"
>
> Analysis: Agreed that the label "ad hominem attack" is not an accurate description of what is written.

Exactly. It's a manipulation, namely "playing the victim". Textbook example.

> The message "I think you should take the time to think about this more" does come across as patronizing because it dismisses the argument made by the other party as poorly thought out [...] Tobias could very well have felt personally attacked there and reacted to that rather than what was literally written.

He did not offer any such argument; he only declared it as "not relevant". Doing so is an expression of disrespect, and I'm asking for the necessary respect in the discussion. Rejecting a dismissal is not a dismissal.

Anyway, I will not comment further on Tobias Diez here. It's clear that you are missing too much context, which takes you too far astray in the analysis.
Let's stick to Gonzalo.

> [...] That's what led to the "disputed ticket" procedure, which has its own flaws.

Indeed, the key flaw: To deny abuse and try to appease abusers.
The effect, very obvious by now: That abusers are encouraged and emboldened.

> Fragment 4 (https://github.com/sagemath/sage/pull/36561#discussion_r1575449665):
> tornaria: "Promote out of beta" doesn't mean it didn't work before 68.1.0, but as I said above, it's not that important. ISTM that 61.0 might be good enough, but again, this is not so important. Cf. <https://setuptools.pypa.io/en/latest/history.html#v61-0-0>
> mkoepe: Surely you are not suggesting that it would be an improvement to use beta-quality features for our production work.
> tornaria: I don't know if you are trying to play some kind of game here. I rather play with my kids. Here I'm trying to help you merge this PR.
> mkoepe: I am not; and this is an unwelcome speculation.
> tornaria: Do you realize I've felt unwelcomed way before your last comment?
>
> Analysis: The PR is about reorganizing metadata and in the process ups the minimum version requirement for setuptools to a relatively new one (68.1.1). Tornaria points out various drawbacks of this [...]

No. Tornaria is not discussing with me. He has no point, and he knows that; he is performing for an uninformed audience, trying to create the impression of a discussion between peers.

> among which that (at the time) several popular linux distributions do not ship such a recent setuptools.

No, this concern is fictional and performative; in technical terms, it's FUD.
- First of all, Gonzalo has no interest in the distribution sage-conf whatsoever.

- These version constraints have already taken effect on various other of our packages ever since the constraints were bumped in <https://github.com/sagemath/sage/pull/36123> (merged October 2023) for the specific purpose of supporting tickets like this one.

- There's no point in having sage-conf support a version of setuptools that our other distributions do not.

- Finally and most importantly: there's simply no packaging scenario in which a Linux distribution with old setuptools can possibly matter -- by the very definition of the build isolation (PEP517/518) that is defined in this file.

- Oh, and `_scope_` (I'll discuss this below).

Contrast: This is what a meaningful discussion of these version

constraints looks like:

https://github.com/3-manifolds/Spherogram/pull/56#discussion_r1495171552

> The reason for the version lower bound given is that with 68.1.0, the required option was declared out of beta. However, as Tornaria comments, the required functionality was available long before. [...] dismissing a compromise on such formal grounds may well look like an overly bureaucratic ploy to avoid making any changes. It is quite probable

... note there's the "probable" again...

> that Tornaria felt he pointed out that the required features had been stable for a long while to you in his earlier comment and hence was irritated that something he considered was a valuable piece of information was dismissed on bureaucratic grounds. From this point it's clear that the discussion will not lead to consensus anyway and essentially Tornaria signs off in desperation with a snarky message,

Well, he for sure shouldn't be doing this because he's already on notice for it.

- If nothing else, it's a public display of disrespect, and as such also part of his dominance posturing strategy.

- "Playing games" is _not_ just a random accusation; rather, it doubles down on

https://github.com/sagemath/sage/pull/37399#discussion_r1527291079 (previously reported).

- Note that on Apr 13 I asked the committee to "communicate to him that a public apology and a change of behavior is necessary." (Has this been done?)

But more specifically: Tornaria knows exactly that setting the version bound as the one in which the feature was declared stable is undeniably "best practices". He continues nevertheless --- because he is writing for an audience, not discussing with me. My "Surely you are not suggesting..." calls him out for this dishonesty and manipulation.

But I have to agree with you, after all, that his "some kind of game" remark is not necessarily a manipulation. It could as well be something else, namely narcissistic rage -- triggered by the realization that his attempted manipulation had just failed.

> [...] It does not look like Tornaria is "playing the victim" in expressing his feeling of being unwelcome. In fact, for a bystander reading through the exchange, Tornaria raises various credible concerns together with fairly simple solutions which are all met with outright dismissal, even though he sketches reasons and incentives for you to engage in reaching a compromise. It does read like his comments were unwelcome from the start.

I agree that this could be a bystander's reading of this passage of the discussion.

Now it's time to consider the missing context here on this PR:

- His comments are outside of the scope of the PR, and he made them

after I have asked multiple times in this PR to respect the scope

(<https://github.com/sagemath/sage/pull/36561#issuecomment-2043211106>,

<https://github.com/sagemath/sage/pull/36561#issuecomment-2041797333>)

-- an important principle of our review code. That he insists on discussing them is already an expression of disrespect.

- Thus he is playing the victim for his disrespectful behavior not being "welcome".

- Ignoring the scope of the PR and making the PR liable for design decisions made in the past (and reviewed by past reviewers...!) is a manipulation by itself.

And the missing larger context from other interactions with Tornaria:

- Once more I have to bring up the "let's not play games" / "do it without tricks" post

(https://github.com/sagemath/sage/pull/37399#discussion_r1527291079).

The committee's explanation from Apr 7: "Gonzalo realized that and struck it out. We think he did a good job in

clarifying that he did not mean to say what he wrote first." Well, no.

People do not make such "mistakes". This is a textbook case of

"testing boundaries": He was trying out how far he can go with public displays of disrespect.

- The entirety of PR "Make sage-conf optional"

(<https://github.com/sagemath/sage/pull/37138>) is an expression of disrespect by its supporters, including Tornaria, which can be summarized as 'We'll just break someone else's design and define that it couldn't possibly matter.'

- The entirety of PR "Remove noisy documentation preview comment"

(<https://github.com/sagemath/sage/pull/37387>) and the attitude there that the development workflow used by the top contributor to Sage can just be casually dismissed as if it could not possibly matter.

- Various attempts to create public doubt about the quality of my work. Posturing to set rules by making "authoritative" declarations about practices and style, e.g., in

<https://github.com/sagemath/sage/pull/37009>. Declarations about my work of years, e.g., "the whole modularization business is just too complicated, period", <https://github.com/sagemath/sage/pull/37178> etc.

> Interestingly, the pull request opens with:

> This is a trivial "chore" PR. It updates Python metadata to the latest format. No controversies about the current format are known about in the Python community. In a typical open source project, someone in a Maintainer role would open a PR and then immediately merge it

> Whereas the comments make clear it turns out to be not so trivial after all: the version requirement bump is a non-trivial change and has real knock-on effects.

No, not at all. That you seem to think that just means that Gonzalo's manipulation has succeeded.

> Designating a ticket as "trivial" can undermine trust and put more friction on future interactions and collaborations.

Not at all, that's an absurd and offensive take.

Designating a ticket as "trivial" gives a valuable signal to reviewers how much time and effort they can expect to spend on its review.

And my designation of this ticket is correct. The review was completed by Francois on Nov 2, 2023, and no meaningful discussion whatsoever has happened after. (Note that the correctness of the designation is also supported by the examples of the PRs with the same content that I submitted to various projects in which the review process and governance function normally.)

Best,
Matthias

--

Dr. Matthias Koeppel <http://www.math.ucdavis.edu/~mkoeppel>
Professor of Mathematics